

Homework 6

References

- Lectures 21-23 (inclusive).

Instructions

- Type your name and email in the "Student details" section below.
- Develop the code and generate the figures you need to solve the problems using this notebook.
- For the answers that require a mathematical proof or derivation you should type them using latex. If you have never written latex before and you find it exceedingly difficult, we will likely accept handwritten solutions.
- The total homework points are 100. Please note that the problems are not weighed equally.

In [126...

```
import numpy as np
np.set_printoptions(precision=3)
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set(rc={"figure.dpi":100, "savefig.dpi":300})
sns.set_context("notebook")
sns.set_style("ticks")

import scipy
import scipy.stats as st
import urllib.request
import os

def download(
    url : str,
    local_filename : str = None
):
    """Download a file from a url.

    Arguments
    url          -- The url we want to download.
    local_filename -- The filename to write on. If not
```

```

        specified
    """
    if local_filename is None:
        local_filename = os.path.basename(url)
        urllib.request.urlretrieve(url, local_filename)

    try:
        import GPpy
    except:
        _=!pip install GPpy
        import GPpy
    import scipy.stats as st

```

Student details

- **First Name:** Benjamin
- **Last Name:** McAteer
- **Email:** bmcateer@purdue.edu

Problem 1 - Defining priors on function spaces

In this problem we are going to explore further how Gaussian processes can be used to define probability measures over function spaces. To this end, assume that there is a 1D function, call it $f(x)$, which we do not know. For simplicity, assume that x takes values in $[0, 1]$. We will employ Gaussian process regression to encode our state of knowledge about $f(x)$ and sample some possibilities for it. For each of the cases below:

- assume that $f \sim \text{GP}(m, k)$ and pick a mean ($m(x)$) and a covariance function $k(x)$ that match the provided information.
- write code that samples a few times (up to five) the values of $f(x)$ at a 100 equidistant points between 0 and 1.

Part A - Super smooth function with known length scale

Assume that you hold the following beliefs

- You know that $f(x)$ has as many derivatives as you want and they are all continuous
- You don't know if $f(x)$ has a specific trend.
- You think that $f(x)$ has "wiggles" that are approximately of size $\Delta x = 0.1$.
- You think that $f(x)$ is between -4 and 4.

Answer:

I am doing this for you so that you have a concrete example of what is requested.

The mean function should be:

$$m(x) = 0.$$

The covariance function should be a squared exponential:

$$k(x, x') = s^2 \exp \left\{ -\frac{(x - x')^2}{2\ell^2} \right\},$$


with variance:

$$s^2 = k(x, x) = \mathbb{V}[f(x)] = 4,$$

and lengthscale $\ell = 0.1$. We chose the variance to be 4 so that with (about) 95% probability the values of $f(x)$ are between -4 and 4.

In [127...

```
#part 1A
# Define the covariance function
k = GPy.kern.RBF(1)
k.lengthscale = 0.1
k.variance = 4.0
# Sample
xs = np.linspace(0, 1, 100)
# The mean function at xs
ms = np.zeros(xs.shape)
# Find the covariance matrix. You need to add a small number
# to the diagonal to ensure numerical stability
nugget = 1e-6
K = k.K(xs[:, None]) + nugget * np.eye(xs.shape[0])
# A multivariate normal that can be used to sample the function values
F = st.multivariate_normal(mean=ms.flatten(), cov=K)
# Take the function samples
f_samples = F.rvs(size=5)
# Plot the samples
fig, ax = plt.subplots()
ax.plot(xs, f_samples.T, 'r', lw=0.5);
plt.title('Part 1A');
```



Part B - Super smooth function with known ultra small length scale

Assume that you hold the following beliefs

- You know that $f(x)$ has as many derivatives as you want and they are all continuous
- You don't know if $f(x)$ has a specific trend.
- You think that $f(x)$ has "wiggles" that are approximately of size $\Delta x = 0.05$.
- You think that $f(x)$ is between -3 and 3.

Answer:

The mean function should be:

$$m(x) = 0.$$

The covariance function should be a squared exponential:

$$k(x, x') = s^2 \exp \left\{ -\frac{(x - x')^2}{2\ell^2} \right\},$$

with variance:

$$s^2 = k(x, x) = \mathbb{V}[f(x)] = 3,$$


and lengthscale $\ell = 0.05$. We chose the variance to be 3 so that with (about) 95% probability the values of $f(x)$ are between -3 and 3.

In [128...

```

###part 1B
# Define the covariance function
k = GPy.kern.RBF(1)
k.lengthscale = 0.05
k.variance = 3.0
# Sample
xs = np.linspace(0, 1, 100)
# The mean function at xs
ms = np.zeros(xs.shape)
# Find the covariance matrix. You need to add a small number
# to the diagonal to ensure numerical stability
nugget = 1e-6
K = k.K(xs[:, None]) + nugget * np.eye(xs.shape[0])
# A multivariate normal that can be used to sample the function values
F = st.multivariate_normal(mean=ms.flatten(), cov=K)
# Take the function samples
f_samples = F.rvs(size=5)
# Plot the samples
fig, ax = plt.subplots()
ax.plot(xs, f_samples.T, 'r', lw=0.5);
plt.title('Part 1B');

```



Part C - Continuous function with known length scale

Assume that you hold the following beliefs

- You know that $f(x)$ is continuous, nowhere differentiable.
- You don't know if $f(x)$ has a specific trend.
- You think that $f(x)$ has "wiggles" that are approximately of size $\Delta x = 0.1$.
- You think that $f(x)$ is between -5 and 5.

Hint: Use `GPy.kern.Exponential` .

Answer:

The mean function should be:

$$m(x) = 0.$$

The covariance function should be an SDE:

$$k(x, x') = \sigma^2 \left(1 + \frac{(x - x')^2}{2\alpha\ell^2}\right)^{-\alpha},$$

as α approaches ∞

with variance:


$$s^2 = k(x, x) = \mathbb{V}[f(x)] = 5,$$

and lengthscale $\ell = 0.1$. We chose the variance to be 5 so that with (about) 95% probability the values of $f(x)$ are between -5 and 5.

In [129...

```
###part 1C
# Define the covariance function
k = GPy.kern.Exponential(1)

k.lengthscale = 0.1
k.variance = 5.0
# Sample
xs = np.linspace(0, 1, 100)
# The mean function at xs
ms = np.zeros(xs.shape)
# Find the covariance matrix. You need to add a small number
# to the diagonal to ensure numerical stability
nugget = 1e-6
K = k.K(xs[:, None]) + nugget * np.eye(xs.shape[0])
# A multivariate normal that can be used to sample the function values
F = st.multivariate_normal(mean=ms.flatten(), cov=K)
# Take the function samples
f_samples = F.rvs(size=5)
# Plot the samples
fig, ax = plt.subplots()
ax.plot(xs, f_samples.T, 'r', lw=0.5);
plt.title('Part 1C');
```



Part D - Smooth periodic function with known length scale

Assume that you hold the following beliefs

- You know that $f(x)$ is smooth.
- You know that $f(x)$ is periodic with period 0.1.
- You don't know if $f(x)$ has a specific trend.
- You think that $f(x)$ has "wiggles" that are approximately of size $\Delta x = 0.5$ of the period.
- You think that $f(x)$ is between -5 and 5.

Hint: Use `GPy.kern.StdPeriodic` .

Answer:

The mean function should be:

$$m(x) = 0.$$

The covariance function should be an Standard Periodic:

$$k(x, x') = \sigma^2 \exp\left(-\frac{2 \sin^2(\pi|x - x'|/p)}{\ell^2}\right)$$

with period:

$$p = 0.1$$

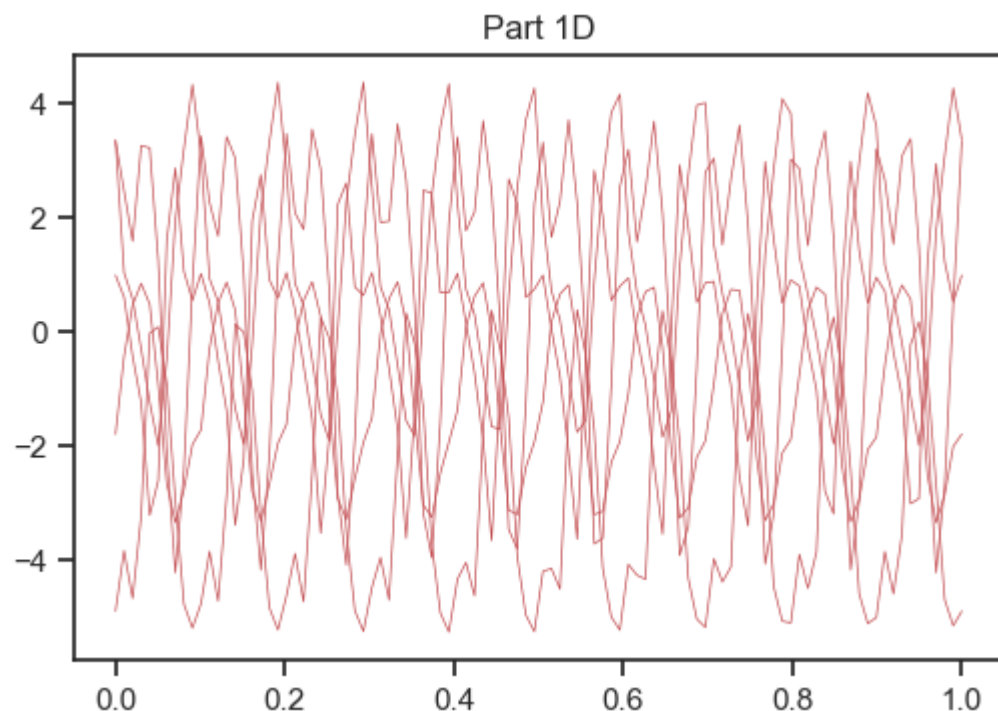
with variance:

$$s^2 = k(x, x) = \mathbb{V}[f(x)] = 5,$$

and lengthscale $\ell = 0.5$. We chose the variance to be 5 so that with (about) 95% probability the values of $f(x)$ are between -5 and 5.

In [130...

```
###part 1D
# Define the covariance function
k = GPy.kern.StdPeriodic(1)
k.period = 0.1
k.lengthscale = 0.5
k.variance = (5.0/1.96)**2
# Sample
xs = np.linspace(0, 1, 100)
# The mean function at xs
ms = np.zeros(xs.shape)
# Find the covariance matrix. You need to add a small number
# to the diagonal to ensure numerical stability
nugget = 1e-6
K = k.K(xs[:, None]) + nugget * np.eye(xs.shape[0])
# A multivariate normal that can be used to sample the function values
F = st.multivariate_normal(mean=ms.flatten(), cov=K)
# Take the function samples
f_samples = F.rvs(size=5)
# Plot the samples
fig, ax = plt.subplots()
ax.plot(xs, f_samples.T, 'r', lw=0.5);
plt.title('Part 1D');
```



Part E - Smooth periodic function with known length scale

Assume that you hold the following beliefs

- You know that $f(x)$ is smooth.
- You know that $f(x)$ is periodic with period 0.1.
- You don't know if $f(x)$ has a specific trend.
- You think that $f(x)$ has "wiggles" that are approximately of size $\Delta x = 0.1$ of the period (**the only thing that is different compared to D**).
- You think that $f(x)$ is between -5 and 5.

Hint: Use `GPy.kern.StdPeriodic` .

Answer: The mean function should be:

$$m(x) = 0.$$

The covariance function should be an Standard Periodic:

$$k(x, x') = \sigma^2 \exp\left(-\frac{2 \sin^2(\pi|x - x'|/p)}{\ell^2}\right)$$

with period:

$$p = 0.1$$


with variance:

$$s^2 = k(x, x) = \mathbb{V}[f(x)] = 5,$$

and lengthscale $\ell = 0.1$. We chose the variance to be 5 so that with (about) 95% probability the values of $f(x)$ are between -5 and 5.

In [131...

```
###part 1E
# Define the covariance function
k = GPy.kern.StdPeriodic(1)
k.period = 0.1
k.lengthscale = 0.1
k.variance = (5.0/1.96)**2
# Sample
xs = np.linspace(0, 1, 100)
# The mean function at xs
ms = np.zeros(xs.shape)
# Find the covariance matrix. You need to add a small number
# to the diagonal to ensure numerical stability
nugget = 1e-6
K = k.K(xs[:, None]) + nugget * np.eye(xs.shape[0])
# A multivariate normal that can be used to sample the function values
F = st.multivariate_normal(mean=ms.flatten(), cov=K)
# Take the function samples
f_samples = F.rvs(size=5)
# Plot the samples
fig, ax = plt.subplots()
ax.plot(xs, f_samples.T, 'r', lw=0.5);
plt.title('Part 1E');
```



Part F - The sum of two functions

Assume that you hold the following beliefs

- You know that $f(x) = f_1(x) + f_2(x)$, where:
 - $f_1(x)$ is smooth with variance 2 and lengthscale 0.5
 - $f_2(x)$ is continuous, nowhere differentiable with variance 0.1 and lengthscale 0.1

Hint: Use must create a new covariance function that is the sum of two other covariances.

In [132...

```
###part 1F
# Define the covariance function
k1 = GPy.kern.StdPeriodic(1)
k1.lengthscale = 0.1
k1.variance = (2.0/1.96)**2


k2 = GPy.kern.Exponential(1)
k2.lengthscale = 0.1
k2.variance = 0.1
```

```

# Sample
xs = np.linspace(0, 1, 100)
# The mean function at xs
ms = np.zeros(xs.shape)
# Find the covariance matrix. You need to add a small number
# to the diagonal to ensure numerical stability
nugget = 1e-6

#k = k1 + k2
#K = k.K(xs[:, None]) + nugget * np.eye(xs.shape[0])
K = k1.K(xs[:, None]) + nugget * np.eye(xs.shape[0]) + k2.K(xs[:, None]) + nugget * np.eye(xs.shape[0])
# A multivariate normal that can be used to sample the function values
F = st.multivariate_normal(mean=ms.flatten(), cov=K)
# Take the function samples
f_samples = F.rvs(size=5)
# Plot the samples
fig, ax = plt.subplots()
ax.plot(xs, f_samples.T, 'r', lw=0.5);
plt.title('Part 1F');

```



Part G - The product of two functions

Assume that you hold the following beliefs


- You know that $f(x) = f_1(x)f_2(x)$, where:
 - $f_1(x)$ is smooth, periodic (period = 0.1), lengthscale 0.1 (relative to the period), and variance 2.
 - $f_2(x)$ is smooth with lengthscale 0.5 and variance 1.

Hint: Use must create a new covariance function that is the product of two other covariances.

In [133...

```
###part 1G
# Define the covariance function
k1 = GPy.kern.StdPeriodic(1)
k1.period = 0.1
k1.lengthscale = 0.1
k1.variance = (2.0/1.96)**2

k2 = GPy.kern.StdPeriodic(1)
k2.lengthscale = 0.5
k2.variance = (2.0/1.96)**2
# Sample
xs = np.linspace(0, 1, 100)
# The mean function at xs
ms = np.zeros(xs.shape)
# Find the covariance matrix. You need to add a small number
# to the diagonal to ensure numerical stability
nugget = 1e-6
K = k1.K(xs[:, None]) + nugget * np.eye(xs.shape[0]) + k2.K(xs[:, None]) + nugget * np.eye(xs.shape[0])
# A multivariate normal that can be used to sample the function values
F = st.multivariate_normal(mean=ms.flatten(), cov=K)
# Take the function samples
f_samples = F.rvs(size=5)
# Plot the samples
fig, ax = plt.subplots()
ax.plot(xs, f_samples.T, 'r', lw=0.5);
plt.title('Part 1G');
```



Problem 2


The National Oceanic and Atmospheric Administration (NOAA) has been measuring the levels of atmospheric CO₂ at the Mauna Loa, Hawaii. The measurements start on March 1958 and go all the way to January 2016. The data can be found [here](#). The Python script below, downloads and plots the data set.

```
In [134... url = "https://github.com/PredictiveScienceLab/data-analytics-se/raw/master/lecturebook/data/mauna_loa_co2.txt"
download(url)
```

```
In [135... data = np.loadtxt('mauna_loa_co2.txt')
```

```
In [136... #load data
t = data[:, 2][:, None] #time (in decimal dates)
y = data[:, 4][:, None] #CO2 level (mole fraction in dry air, micromol/mol, abbreviated as ppm)
fig, ax = plt.subplots(1, figsize = (15, 10), dpi=100)
ax.plot(t, y, '.')
```

```
ax.set_xlabel('$t$ (year)')  
ax.set_ylabel('$y$ (CO2 level in ppm)');
```



Overall, we observe a steady growth of CO2 levels. The wiggles correspond to seasonal changes. Since the vast majority of the population inhabits the Northern hemisphere, fuel consumption goes up during the Northern winters and CO2 emissions follow. Our goal is to study this dataset with Gaussian process regression. Specifically we would like to predict the evolution of the CO2 levels from Feb 2018 to Feb 2028 and quantify our uncertainty about this prediction.

It's always a good idea to work with a scaled version of the inputs and the outputs. We are going to scale the times as follows:

$$t_s = t - t_{\min}.$$


So, time is still in fractional years, but we start counting at zero instead of 1950. We scale the y 's as:

$$y_s = \frac{y - y_{\min}}{y_{\max} - y_{\min}}.$$

This takes all the y between 0 and 1. Here is how the scaled data look like:

In [137...

```
t_s = t - t.min()
y_s = (y - y.min()) / (y.max() - y.min())
fig, ax = plt.subplots(1, figsize = (15, 10), dpi=100)
ax.plot(t_s, y_s, '.')
ax.set_xlabel('$t_s$ (Scaled year)')
ax.set_ylabel('$y_s$ (Scaled CO2 level)$');
```



In what follows, just work with the scaled data as you develop your model. Scale back to the original units for your final predictions.

Part A - Naive approach

Use a zero mean Gaussian process with a squared exponential covariance function to fit the data and make the required prediction (ten years after the last observation).

Answer:


In [138...

```
##Part A
k= GPy.kern.RBF(1)
naive_model = GPy.models.GPRegression(t_s, y_s, k)
naive_model.optimize();
```

Predict everything:

In [139...

```
tss = np.linspace(0, t_s.max() + 10, 200)[: , None]
ys, vs = naive_model.predict(tss)
ls = ys - 1.96 * np.sqrt(vs)
us = ys + 1.96 * np.sqrt(vs)
fig, ax = plt.subplots(1, figsize = (15, 10), dpi=100)
ax.plot(tss, ys, color='blue', label='Posterior mean')
ax.fill_between(tss.flatten(), ls.flatten(), us.flatten(), color='blue', alpha=0.25)
ax.plot(t_s, y_s, '.', label='Scaled observed data')
ax.set_xlabel('$t_s$ (Scaled year)')
ax.set_ylabel('$y_s$ (Scaled CO2 level)');
```



Notice that the squared exponential covariance captures the long terms, but it fails to capture the seasonal fluctuations. As a matter of fact the seasonal fluctuations are treated as noise. This is clearly false. How can we fix it?

Part B - Improving the prior covariance

Now use the ideas of Problem 1, to come up with a covariance function that exhibits the following characteristics clearly visible in the data (call $f(x)$ the scaled CO2 level).

- $f(x)$ is smooth
- $f(x)$ has a clear trend with a multi-year lengthscale (it is also an increasing trend, but we are not going to impose this)
- $f(x)$ has seasonal fluctuations with a period of one year
- $f(x)$ exhibits small fluctuations within its period.

Use summation and multiplication of simple covariance functions to create a covariance function that exhibits these trends. Sample a few times from it.

Hint: Do not attempt to fit the data in any way. Just try to find a covariance function that has the right features. We also do not care about getting the parameters 100% right at this point. The parameters will be optimized later.

Answer:

In [140...

```
k1= GPy.kern.RBF(1) #short trend
k1.lengthscale= 1/len(y_s)

k2 = GPy.kern.StdPeriodic(1) #seasonal
k2.period = 1
k2.lengthscale = 1

k3 = GPy.kern.RBF(1)#Long trend
k3.lengthscale= len(y_s)

k = (k1 + k2) * k3


Improved_model = GPy.models.GPRegression(t_s, y_s, k)
Improved_model.optimize()

tss = np.linspace(0, t_s.max() + 10, 200)[: , None]
ys, vs = Improved_model.predict(tss)
ls = ys - 1.96 * np.sqrt(vs)
us = ys + 1.96 * np.sqrt(vs)
fig, ax = plt.subplots(1, figsize = (15, 10), dpi=100)
ax.plot(tss, ys, color='blue', label='Posterior mean')
ax.fill_between(tss.flatten(), ls.flatten(), us.flatten(), color='blue', alpha=0.25)
ax.plot(t_s, y_s, '.', label='Scaled observed data')
ax.set_xlabel('$t_s$ (Scaled year)')
```

```
ax.set_ylabel('$y_s$ (Scaled CO2 level)');  
plt.title('Problem 2B')
```

D:\School\Programing\Anaconda\envs\BME301\lib\site-packages\paramz\transformations.py:111: RuntimeWarning:overflow encountered in expm1

Out[140... Text(0.5, 1.0, 'Problem 2B')



Part C - Predicting the future


Use a zero mean Gaussian process with the covariance function you picked above to do Gaussian process regression and make the required prediction (ten years after the last observation).

Answer:

In [141...

```
##part C
fig, ax = plt.subplots(1, figsize = (15, 10), dpi=100)
ax.plot(tss, ys, color='blue', label='Posterior mean')
ax.fill_between(tss.flatten(), ls.flatten(), us.flatten(), color='blue', alpha=0.25)
ax.plot(t_s, y_s, '.', label='Scaled observed data')
ax.set_xlabel('$t_s$ (Scaled year)')
ax.set_ylabel('$y_s$ (Scaled CO2 level)');
plt.title('Problem 2C: Predicting the Future')
plt.xlim(60,70)
plt.ylim(0.8, 1.3)
```

Out[141... (0.8, 1.3)



Part D - Bayesian information criterion

As we have seen in earlier lectures, the Bayesian information criterion (BIC), see [this](#), can be used to compare two models. The criterion says that one should:

- fit the models with maximum likelihood,
- and compute the quantity:

$$\text{BIC} = d \ln(n) - 2 \ln(\hat{L}),$$

where d is the number of model parameters, and \hat{L} the maximum likelihood.

- pick the model with the smallest BIC.

Use BIC to show that the model you constructed in Part C is indeed better than the naïve model of Part A.

Hint: Do a `help(GPy.models.GPRegression)` and you will find a way to get both the number of parameters and the log likelihood. Ask on piazza if you can't find it - or Google it.

Answer:

In [142...

#Part D

```
naive_BIC = -2 * naive_model.log_likelihood() + naive_model._size_transformed() * np.log(len(naive_model.X))
improved_BIC = -2 * Improved_model.log_likelihood() + Improved_model._size_transformed() * np.log(len(Improved_model.X))

print('The Naive model BIC is:', naive_BIC)
print('The Part B&C model BIC is:', improved_BIC)
print('Since the Model from part B&C has a more negative BIC, it is indeed a better model than the part A model')
```

The Naive model BIC is: -3363.207769943941

The Part B&C model BIC is: -4464.4176418006155

Since the Model from part B&C has a more negative BIC, it is indeed a better model than the part A model

Part E - Plot samples from the posterior Gaussian process

Using the model of Part C, plot 5 samples from the posterior Gaussian process between 2018 and 2028.


Hint: You need to use `GPy.models.GPRegression.posterior_samples_f`.

Answer:

In [143...

##Part E

```
f_postsamples = Improved_model.posterior_samples_f(tss,5)
# Plot the samples
fig, ax = plt.subplots(1, figsize = (15, 10), dpi=100)
ax.plot(tss, ys, color='blue', label='Posterior mean')
ax.fill_between(tss.flatten(), ls.flatten(), us.flatten(), color='blue', alpha=0.25)
#I HAVE NO CLUE WHY THE FOR LOOP DIDNT WORK
plt.plot(tss[:,:],f_postsamples[:,0],label = 'Sample 1')
plt.plot(tss[:,:],f_postsamples[:,1],label = 'Sample 2')
plt.plot(tss[:,:],f_postsamples[:,2],label = 'Sample 3')
plt.plot(tss[:,:],f_postsamples[:,3],label = 'Sample 4')
plt.plot(tss[:,:],f_postsamples[:,4],label = 'Sample 5')
ax.set_xlabel('$t_s$ (Scaled year)')
ax.set_ylabel('$y_s$ (Scaled CO2 level)');
plt.title('Problem 2E Posterior Samples zoomed to 2018-2028')
plt.xlim(60,70);
plt.ylim(0.8, 1.3);
plt.legend();
```



Problem 3 - Using Bayesian Global optimization to calibrate an expensive physical model

This is Example 3.1 of [\(Tsilifis, 2014\)](#).

Consider the catalytic conversion of nitrate (NO_3^-) to nitrogen (N_2) and other by-products by electrochemical means. The mechanism that is followed is complex and not well understood. The experiment of (Katsounaros, 2012) confirmed the production of nitrogen (N_2), ammonia (NH_3), and nitrous oxide (N_2O) as final products of the reaction, as well as the intermediate production of nitrite (NO_2^-). The data are reproduced in [Comma-separated values](#) (CSV) and stored in [catalysis.csv](#). The time is measured in minutes and the concentrations are measured in $\text{mmol} \cdot \text{L}^{-1}$. Let's load the data into this notebook using the [Pandas](#) Python module:


```
In [144... url = "https://github.com/PredictiveScienceLab/data-analytics-se/raw/master/lecturebook/data/catalysis.csv"
download(url)
```

```
In [145... # Load the data
import pandas as pd
catalysis_data = pd.read_csv('catalysis.csv')
catalysis_data
```

```
Out[145...
```

	Time	NO3	NO2	N2	NH3	N2O
0	0	500.00	0.00	0.00	0.00	0.00
1	30	250.95	107.32	18.51	3.33	4.98
2	60	123.66	132.33	74.85	7.34	20.14
3	90	84.47	98.81	166.19	13.14	42.10
4	120	30.24	38.74	249.78	19.54	55.98
5	150	27.94	10.42	292.32	24.07	60.65
6	180	13.54	6.11	309.50	27.26	62.54

```
In [146... catalysis_data.plot(style='s', x=0);
```



The theory of catalytic reactions guarantees that the total mass must be conserved. However, this is not the case in our dataset:

```
In [147... catalysis_data.sum(axis=1)
```

```
Out[147... 0    500.00
1    415.09
2    418.32
3    494.71
4    514.28
5    565.40
6    598.95
dtype: float64
```

This inconsistency suggests the existence of an intermediate unobserved reaction product X. ([Katsounaros, 2012](#)) suggested that the following reaction path shown in the following figure.

The dynamical system associated with the reaction is:

$$\begin{aligned}
\frac{d[\text{NO}_3^-]}{dt} &= -k_1 [\text{NO}_3^-], \\
\frac{d[\text{NO}_2^-]}{dt} &= k_1 [\text{NO}_3^-] - (k_2 + k_4 + k_5)[\text{NO}_2^-], \\
\frac{d[\text{X}]}{dt} &= k_2 [\text{NO}_2^-] - k_3[\text{X}], \\
\frac{d[\text{N}_2]}{dt} &= k_3 [\text{X}], \\
\frac{d[\text{NH}_3]}{dt} &= k_4 [\text{NO}_2^-], \\
\frac{d[\text{N}_2\text{O}]}{dt} &= k_5 [\text{NO}_2^-],
\end{aligned}$$

where $[\cdot]$ denotes the concentration of a quantity, and $k_i > 0$, $i = 1, \dots, 5$ are the *kinetic rate constants*.

In this problem, I am going to guide you through the calibration of the parameters of this model so that we match the observations. These problems are also known as *inverse problems*. The problem can, and should, be formulated in a Bayesian way. However, in this homework problem we are going to do it using a classical loss-minimization approach. We will discuss the Bayesian approach for calibrating the same model in a later lecture.

Before you proceed, please read a little bit about the "classical theory of inverse problems:"

Classical theory of inverse problems

Suppose that you have a model (any model really) that predicts a quantity of interest. Let's assume that this model has parameters that you do not know. These parameters could be simple scalars (mass, spring constant, dumping coefficients, etc.) or it could be also be functions (initial conditions, boundary values, spatially distributed constitutive relations, etc.) Let's denote all these parameters with the vector x . Assume that:

$$x \in \mathcal{X} \subset \mathbb{R}^d.$$

Now, let's say we perform an experiment that measures a *noisy* vector:

$$y \in \mathcal{Y} \subset \mathbb{R}^m.$$

Assume that, you can use your model *model* to predict y . It does not matter how complicated your model is. It could be a system of ordinary differential or partial differential equations, or something more complicated. If it predicts y , you can always think of it as a function

from the unknown parameter space \mathcal{X} to the space of y 's, $\mathcal{Y} \subset \mathbb{R}^m$. That is, you can think of it as giving rise to a function:

$$f : \mathcal{X} \rightarrow \mathcal{Y}.$$

The **inverse problem**, otherwise known as the **model calibration** problem is to find the best $x \in \mathcal{X}$ so that:

$$f(x) \approx y.$$

Formulation of Inverse Problems as Optimization Problems

Saying that $f(x) \approx y$ is not an exact mathematical statement. What does it really mean for $f(x)$ to be close to y ? To quantify this, let us introduce a *loss metric*:

$$\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R},$$

such that $\ell(f(x), y)$ is how much our prediction is off if we chose the input $x \in \mathcal{X}$. Equipped with this loss metric, we can formulate the mathematical problem as:

$$\min_{x \in \mathcal{X}} \ell(f(x), y).$$

The Square Loss

The choice of the metric is somewhat subjective (it depends on what it means to be wrong in your problem). However, a very common assumption is that to take the *square loss*:

$$\ell(f(x), y) = \frac{1}{2} \|f(x) - y\|_2^2 = \frac{1}{2} \sum_{i=1}^m (f_i(x) - y_i)^2.$$

For this case, the inverse problem can be formulated as:

$$\min_{x \in \mathcal{X}} \frac{1}{2} \|f(x) - y\|_2^2.$$

Solution Methodologies

We basically have to solve an optimization problem. For the square loss function, if $f(x)$ is linear, then you get the classic least squares problem which has a known solution. Otherwise, you get what is known as *generalized least squares*. There are many algorithms that you could use this problem. Several are implemented in [scipy.optimize](#). If you are able to implement your model as a simple python function,

then you can use them. Alternatively, and this is what we are going to do here, we could use Bayesian global optimization instead. The absolutely, essential thing that you need to provide to these methods is the function they are optimizing, i.e.,

$$L(x, y) = \ell(f(x), y).$$

Back to the catalysis model

Let's now formulate the calibration problem for the catalysis model. We proceed in several steps.

Step 1: Making our life easier by simplifying the notation

Note that this is actually a linear system. To simplify our notation, let's define:

$$\begin{aligned} z_1 &:= [\text{NO}_3^-], \\ z_2 &:= [\text{NO}_2^-], \\ z_3 &:= [\text{X}], \\ z_4 &:= [\text{N}_2], \\ z_5 &:= [\text{NH}_3], \\ z_6 &:= [\text{N}_2\text{O}], \end{aligned}$$

the vector:

$$z = (z_1, z_2, z_3, z_4, z_5, z_6),$$

and the matrix:

$$A(k_1, \dots, k_5) = \begin{pmatrix} -k_1 & 0 & 0 & 0 & 0 & 0 \\ k_1 & -(k_2 + k_4 + k_5) & 0 & 0 & 0 & 0 \\ 0 & k_2 & -k_3 & 0 & 0 & 0 \\ 0 & 0 & k_3 & 0 & 0 & 0 \\ 0 & k_4 & 0 & 0 & 0 & 0 \\ 0 & k_5 & 0 & 0 & 0 & 0 \end{pmatrix} \in \mathbb{R}^{6 \times 6}.$$

With these definitions, the dynamical system becomes:

$$\dot{z} = A(k_1, \dots, k_5)z,$$

with initial conditions

$$z(0) = z_0 = (500, 0, 0, 0, 0, 0) \in \mathbb{R}^6,$$

read directly from the experimental data. What we are definitely going to need is a solver for this system. That's easy. Let's denote the solution of the system at time t by:

$$z(t; k_1, \dots, k_5).$$

Step 2: Scale the unknown parameters to your best of your abilities

The constraints you have on your parameters, the better. If you do have constraints, you would have to use constrained optimization algorithms. The way you scale things depend on the problem. Here we would think as follows:

- k_i has units of inverse time. It is properly appropriate to scale it with the total time which is 180 minutes. So, let's just multiply k_i with 180. This makes the resulting variable dimensionless:

$$\hat{x}_i = 180k_i.$$

- k_i is positive, therefore \hat{x}_i must be positive. So, let's just work with the logarithm of \hat{x}_i :

$$x_i = \log \hat{x}_i = \log 180k_i.$$

- define the parameter vector:

$$x = (x_1, \dots, x_5) \in \mathcal{X} = \mathbb{R}^5.$$

From now on, we will write

$$A = A(x),$$

for the matrix of the dynamical system, and

$$z = z(t; x),$$

for the solution at t given that the parameters are x .

Step 3: Making the connection between our model and the experimental measurements

Our experimental data include measurements of everything except z_3 at times six (6) time instants:

$$t_j = 30j \text{ minutes,}$$

$$j = 1, \dots, 6.$$

Now, let $Y \in \mathbb{R}^{5 \times 6}$ be the experimental measurements:

In [148...

```
catalysis_data[1:]
```

Out[148...

	Time	NO3	NO2	N2	NH3	N2O
1	30	250.95	107.32	18.51	3.33	4.98
2	60	123.66	132.33	74.85	7.34	20.14
3	90	84.47	98.81	166.19	13.14	42.10
4	120	30.24	38.74	249.78	19.54	55.98
5	150	27.94	10.42	292.32	24.07	60.65
6	180	13.54	6.11	309.50	27.26	62.54

You can think of the measurements as vector by flattening the matrix:

$$y = \text{vec}(Y) \in \mathbb{R}^{30}.$$

Note that `vec` is the vectorization operator.

What is the connection between the solution of the dynamical system $z(t, x)$ and the experimental data? It is as follows:

$$\begin{aligned} z_1(30j; x) &\longrightarrow Y_{j1}, \\ z_2(30j; x) &\longrightarrow Y_{j2}, \\ z_4(30j; x) &\longrightarrow Y_{j3}, \\ z_5(30j; x) &\longrightarrow Y_{j4}, \\ z_6(30j; x) &\longrightarrow Y_{j5}, \end{aligned}$$

for $j = 1, \dots, 6$.

We are now ready to define a function:

$$f : \mathcal{X} \rightarrow \mathcal{Y} = \mathbb{R}_+^{30},$$

as follows:

- Define the matrix function:

$$F : \mathcal{X} \rightarrow \mathbb{R}^{5 \times 6},$$

by:

$$\begin{aligned} F_{j1}(x) &= z_1(30j; x) \longrightarrow Y_{j1}, \\ F_{j2}(x) &= z_2(30j; x) \longrightarrow Y_{j2}, \\ F_{j3}(x) &= z_4(30j; x) \longrightarrow Y_{j3}, \\ F_{j4}(x) &= z_5(30j; x) \longrightarrow Y_{j4}, \\ F_{j5}(x) &= z_6(30j; x) \longrightarrow Y_{j5}, \end{aligned}$$

- And flatten that function:

$$f(x) = \text{vec}(F(x)) \in \mathbb{R}^{30}.$$

Now, we have made the connection with our theoretical formulation of inverse problems crystal clear.

Step 4: Programming our ODE solver and the loss function

In [149...

```
import scipy.integrate

def A(x):
    """
    Return the matrix of the dynamical system.
    """
    # Scale back to the k's
    k = np.exp(x) / 180.
    res = np.zeros((6,6))
    res[0, 0] = -k[0]
    res[1, 0] = k[0]
    res[1, 1] = -(k[1] + k[3] + k[4])
    res[2, 1] = k[1]
    res[2, 2] = -k[2]
    res[3, 2] = k[2]
    res[4, 1] = k[3]
    res[5, 1] = k[4]
    return res
```

```

def g(z, t, x):
    """
    The right hand side of the dynamical system.
    """
    return np.dot(A(x), z)

# The initial conditions
z0 = np.array([500., 0., 0., 0., 0., 0.])

# The times at which we need the solution (experimental times)
t_exp = np.array([30. * j for j in range(1, 7)])

# The experimental data as a matrix
Y = catalysis_data[1:].values[:, 1:]

# The experimental as a vector
y = Y.flatten()

# The full solution of the dynamical system
def Z(x, t):
    """
    Returns the solution for parameters x at times t.
    """
    return scipy.integrate.odeint(g, z0, t, args=(x,))

# The matrix function F (matches to Y)
def F(x, t):
    res = Z(x, t)
    return np.hstack([res[:, :2], res[:, 3:]])

# The function f (matches to y)
def f(x, t):
    return F(x, t).flatten()

# Finally, the loss function that we need to minimize over x:
def L(x, t, y):
    return 0.5 * np.sum((f(x, t) / 500. - y / 500.) ** 2) # We scale for numerical stability

```

Step 5: Minimize the loss function

Let's optimize with scipy.optimize:

In [150...

```
import scipy.optimize

# Initial guess for x
x0 = -2.0 + 2.0 * np.random.rand(5)

# Optimize
res = scipy.optimize.minimize(L, x0, args=(t_exp, y))

print(res)


fun: 0.16530885803340029
hess_inv: array([[ 13.219,  22.107, -58.737, -82.055,  16.42 ],
 [ 22.107,  61.581, -141.608, -192.154,  39.77 ],
 [ -58.737, -141.608,  365.412,  518.833, -91.109],
 [ -82.055, -192.154,  518.833,  795.313, -120.025],
 [  16.42 ,  39.77 , -91.109, -120.025,  32.282]])
jac: array([-6.801e-06,  1.604e-06,  1.984e-06,  9.835e-07,  3.820e-06])
message: 'Optimization terminated successfully.'
nfev: 342
nit: 41
njev: 57
status: 0
success: True
x: array([ 1.926,  2.019,  1.425, -0.526,  0.265])
```


And here is how you can visualize the model with the "best" parameters:

In [151...

```
x = res.x
t = np.linspace(0, 180, 100)
x1 = np.array([1.359, 1.657, 1.347, -.16, -1.01])
Yp = Z(x, t)

fig, ax = plt.subplots(figsize=(10, 10))
catalysis_data.plot(ax=ax, style='s', x=0)
ax.plot(t, Yp[:, 0], color=sns.color_palette()[0], label='Model NO3-')
ax.plot(t, Yp[:, 1], color=sns.color_palette()[1], label='Model NO2-')
ax.plot(t, Yp[:, 2], color=sns.color_palette()[5], label='Model X')
ax.plot(t, Yp[:, 3], color=sns.color_palette()[2], label='Model N2')
ax.plot(t, Yp[:, 4], color=sns.color_palette()[3], label='Model NH3')
ax.plot(t, Yp[:, 5], color=sns.color_palette()[4], label='Model N2O')
plt.legend();
```





Note that the code above will not work every time... Some times it will work and sometimes it won't work. Run it 3-4 times if it accidentally works. There are several problems. Here are the three most relevant in our context:

- `scipy.optimize` needs the gradient of the loss function. Since we do not provide it, it tries to get it using numerical differentiation. Numerical differentiation introduces errors...
- `scipy.optimize` find a local minimum of the loss function. This may actually be a bad local minimum.
- okay, this particular model is not very computationally expensive. But imagine trying to calibrate a model that takes a while for a single evaluation (e.g., a finite element model). Then, using `scipy.optimize` (especially without supplying the derivatives), is doomed to fail.

To overcome these difficulties, you have to use Bayesian global optimization to solve the problem. Note that in the hands-on activities, we introduced this code:

In [152...

```
def ei(m, sigma, ymax, psi=0.):
    u = (m - ymax) / sigma
    ei = sigma * (u * st.norm.cdf(u) + st.norm.pdf(u))
    ei[sigma <= 0.] = 0.
    return ei

def maximize(f, gpr, domain, num_candidates=10000,
            alpha=ei, psi=0., max_it=6):
    """
    Optimize f using a limited number of evaluations.

    :param f:      The function to optimize.
    :param gpr:     A Gaussian process model to use for representing our state of knowledge.
    :param X_design: The set of candidate points for identifying the maximum.
    :param alpha:   The acquisition function.
    :param psi:     The parameter value for the acquisition function (not used for EI).
    :param max_it:  The maximum number of iterations.
    """
    af_all = []
    print('Iteration\tCurrent best objective \tCurrent acquisition func. value')
    dim = gpr.X.shape[1]
    for count in range(max_it):
        X_design = domain[:, 0] + \
            (domain[:, 1] - domain[:, 0]) * \
```

```

        np.random.rand(num_candidates, dim)
    m, sigma2 = gpr.predict(X_design)
    sigma = np.sqrt(sigma2)
    af_values = alpha(m, sigma, gpr.Y.max(), psi=psi)
    i = np.argmax(af_values)
    X = np.vstack([gpr.X, X_design[i:(i+1), :]])
    y = np.vstack([gpr.Y, [f(X_design[i, :])]])
    gpr.set_XY(X, y)
    # Uncomment the following to optimize the hyper-parameters
    #gpr.optimize()
    idx_opt = np.argmax(gpr.Y.flatten())
    f_opt = gpr.Y[idx_opt, 0]
    print('{0:d}\t\t{1:1.2f}\t\t{2:1.2f}'.format(count + 1, f_opt, af_values[i, 0]))
    x_opt = np.array(gpr.X[idx_opt])
    return x_opt, f_opt, gpr

```

The code *maximizes* a function, but you want to *minimize* the loss. To recast the problem as a maximization problem, you need to work with *minus the loss*. Also, the code does not allow for a function with extra parameters (like the `t_exp` and the `y` we have for `L`). Here is the function that you should be optimizing:

```
In [153... h = lambda x: -L(x, t_exp, y)
```

Part A - Perform multivariate Gaussian process regression on an initial set of data

We are going to search for the best parameters x within the set $[-2, 2]^5$. Consider the following two datasets consisting of parameter and minus loss pairs:

```
In [154... # Initial training points
n_init_train = 100
X_init_train = -2.0 + 4.0 * np.random.rand(n_init_train, 5)
Y_init_train = np.array([h(x) for x in X_init_train])[:, None]
```

Use a squared exponential covariance function with automatic relevance determination to do Gaussian process regression with `X_init_train` and `Y_init_train`.

Hint: You may want to experiment by constraining the likelihood noise of your model to be very small, say 10^{-6} . This is because the observations of the loss do not really have any noise.

Answer:

In [155...

```

##part A
# Define the covariance function
k = GPy.kern.RBF(5, ARD = True)
#k.lengthscale = 1
GP= GPy.models.GPRegression(X_init_train,Y_init_train,k)
GP.optimize(messages=True);

```

Running L-BFGS-B (Scipy implementation) Code:

runtime	i	f	g
00s00	0000	1.543247e+02	nan
00s01	0003	3.646967e+01	8.512773e+02
00s05	0014	-9.419900e+01	1.417629e+01
00s09	0025	-9.946722e+01	5.488522e-03
00s10	0029	-9.946966e+01	2.460373e-03
00s13	0037	-9.946995e+01	7.670067e-07
00s14	0039	-9.946995e+01	3.623508e-08

Runtime: 00s14

Optimization status: Converged

Part B - Inspecting your model

Use the lengthscale information to rank the model parameters according their effect on the calibration loss.

Answer:

In [156...

```

###part B
print(GP.kern.lengthscale)
print(' The larger lengthscales are less important, therefore index 0,3,4,1, and 2 are the rankings of most important to

```

index	GP_regression.rbf.lengthscale	constraints	priors
[0]	2.629	+ve	
[1]	4.948	+ve	
[2]	16.297	+ve	
[3]	5.636	+ve	
[4]	6.973	+ve	

The larger lengthscales are less important, therefore index 0,3,4,1, and 2 are the rankings of most important to least important respectively

Part C - Diagnostics

Here are some test data:

```
In [157...
# Test points
n_test = 50
X_test = -2.0 + 4.0 * np.random.rand(n_test, 5)
Y_test = np.array([h(x) for x in X_test][:, None])
```

Do the following:

- Predictions vs observations plot
- Standarized errors plot

Answer:

```
In [158...
#part C
from sklearn.metrics import mean_squared_error
# Test points
n_test = 50
X_test = -2.0 + 4.0 * np.random.rand(n_test, 5)
Y_test = np.array([h(x) for x in X_test][:, None])

Y_test_m, Y_test_v = GP.predict(X_test, full_cov=False)
error_m = np.sqrt( mean_squared_error(Y_test, Y_test_m) )

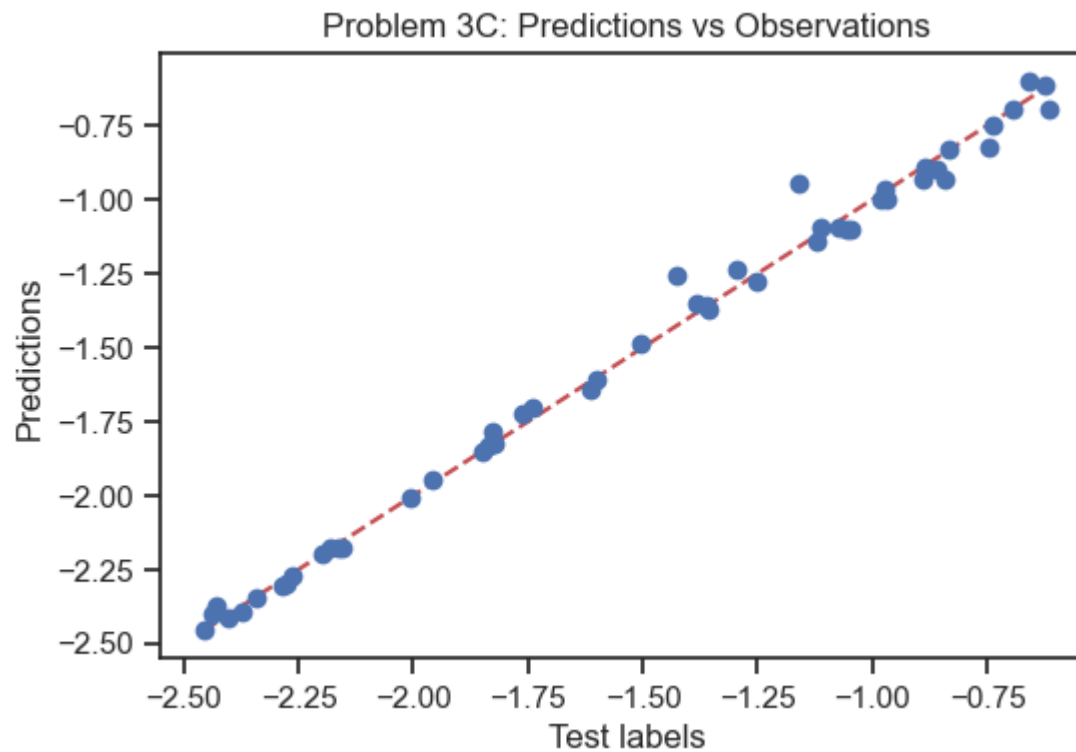
fig, ax = plt.subplots()
ys = np.linspace( np.min(Y_test), np.max(Y_test), 100 )
ax.plot(ys, ys, 'r--', label='$x=y$')
ax.plot(Y_test, Y_test_m, 'o')
ax.set_title('Predictions vs True test values')
ax.set_xlabel('Test labels')
ax.set_ylabel('Predictions');
plt.title('Problem 3C: Predictions vs Observations')


eps = (Y_test_m - Y_test) / np.sqrt(Y_test_v)
idx = np.arange(1,eps.shape[0] + 1)

fig, ax = plt.subplots()
ax.plot(idx, eps, 'o', label='Standarized errors')
```

```
ax.plot(idx, 1.96 * np.ones(eps.shape[0]), 'r--')
ax.plot(idx, -1.96 * np.ones(eps.shape[0]), 'r--')
ax.set_xlabel('$i$')
ax.set_ylabel('$\epsilon_i$');
plt.title('Problem 3C: Standardized Errors')
```

Out[158...] Text(0.5, 1.0, 'Problem 3C: Standardized Errors')





Part D - Calibrate the model with Bayesian global optimization

Now use Bayesian global optimization with expected improvement to calibrate your model using the GP that you built above as the starting point. Do not expect this to give you a perfect model. But it will be better than nothing. We will get the best possible model in the next homework assignment.

Hint: Here you basically need to read the docstring of `maximize` and use it correctly.

In [159...

```
# For your convenience, the `domain` argument of minimize should be:
domain = np.array([[-2, 2], [-2, 2], [-2, 2], [-2, 2], [-2, 2]])
# Run maximize here:
# Your code here
x_opt, f_opt, GPr = maximize(h, GP, domain, max_it = 30) # The best parameters
```

Iteration	Current best objective	Current acquisition func. value
1	-0.21	0.28
2	-0.21	0.23

3	-0.20	0.05
4	-0.18	0.06
5	-0.18	0.01
6	-0.17	0.05
7	-0.17	0.01
8	-0.17	0.02
9	-0.17	0.03
10	-0.17	0.02
11	-0.17	0.02
12	-0.17	0.01
13	-0.17	0.02
14	-0.17	0.01
15	-0.17	0.01
16	-0.17	0.00
17	-0.17	0.00
18	-0.17	0.02
19	-0.17	0.04
20	-0.17	0.01
21	-0.17	0.01
22	-0.17	0.02
23	-0.17	0.02
24	-0.17	0.00
25	-0.17	0.02
26	-0.17	0.00
27	-0.17	0.04
28	-0.17	0.01
29	-0.17	0.00
30	-0.17	0.02

Use this code to plot your calibrated model:

In [160...

```
x = x_opt
t = np.linspace(0, 180, 100)
x1 = np.array([1.359, 1.657, 1.347, -.16, -1.01])
Yp = Z(x, t)

fig, ax = plt.subplots(figsize=(10, 10))
catalysis_data.plot(ax=ax, style='s', x=0)
ax.plot(t, Yp[:, 0], color=sns.color_palette()[0], label='Model NO3-')
ax.plot(t, Yp[:, 1], color=sns.color_palette()[1], label='Model NO2-')
ax.plot(t, Yp[:, 2], color=sns.color_palette()[5], label='Model X')
ax.plot(t, Yp[:, 3], color=sns.color_palette()[2], label='Model N2')
ax.plot(t, Yp[:, 4], color=sns.color_palette()[3], label='Model NH3')
ax.plot(t, Yp[:, 5], color=sns.color_palette()[4], label='Model N2O')
plt.legend();
```

